

# Report Title

Author Name1\* and Author Name2\*\*

\*Dept. of Information Studies

\*\*Dept. of Physics & Astronomy

University College London, WC1E 6BT

*{email1,email2}@ucl.ac.uk*

April 30, 2020

## Abstract

The abstract text goes here. Here is a bit more of my abstract.

## 1 Introduction

Here is the text of your introduction.

Here is a simple equation

$$\alpha = \sqrt{\beta} \tag{1}$$

I can refer to my equation by saying look at Equation (1).

Here is a multi-line equation:

$$\begin{aligned} f(\mathbf{x}; \mathbf{w}) &= \sigma(\mathbf{w}^T \mathbf{x}) \\ &= \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} \end{aligned} \tag{2}$$

I can point people to literature like this [3] (a conference paper) or this [2] (a journal paper), or even this [1, Sec. 4.1] (a book, with section indicated). See the file `report.bib` for how these bibliographic entries are defined.

It helps to have a convention to the font used for particular mathematical objects. A common convention is: for simple values use lower case latin characters  $s$  (this might include observed values in probabilistic work); for sets use caligraphic upper case  $\mathcal{S}$ ; for vectors of values use bold latin  $\mathbf{s}$ ; for matrices use upper case latin  $S$ ; and for parameters (including parameters in probabilistic models) use lower case greek letters (scalar),  $\sigma$ , bold greek (vectors),  $\boldsymbol{\sigma}$ , and possibly upper-case greek for matrices of parameters  $\Sigma$ . However, it is up to you to develop your own style.

### 1.1 Subsection Heading

Write your subsection text here. Here I am referencing Section 1. Figures and tables can be inserted and referred to, e.g. Figure 1 and Table 1.

	Interested	Understand
Linear Regression	N	Y
Neural Networks	Y	N

Table 1: Some results as a table.



Figure 1: Simulation Results

## 1.2 Quoting and Emphasis

When you quote or emphasise things, you have a few options. I tend to put things in italics, e.g. Luke told me, *Put it in italics!*. If you prefer to put it in quote-marks then do so, but make sure you use left and right quotes properly, e.g. Luke told me, “Use left quotes to open, and right quotes to close!”.

## 2 Background

Your background section text goes here...

## 3 Method

### 3.1 Game Mechanics

We have decided to follow the exact game mechanics in the actual game with no reduction and simplification in the rules. This means that we will always start with a random board with 2 tiles with values on each of the tiles "2" or "4". Then, after each valid move, a new tile with values "2" or "4" will be generated on a random empty tile. We loose when no valid moves are allowed.

### 3.2 Representation and Q-value Calculation

In our experiment, we have tried 2 types of representation, which are the simple representation and the relationship representation. They will be explained as follows. To calculate the q-value, we will use the linear approximation approach for simplicity.

#### 3.2.1 Simple Representation

This is the simplest type of representation of our game, which is just simply a vector of 16 elements with each element representing the values of the individual tiles. For example, if we look at figure 2, the state will simply be:

$$\phi = (0, 0, 2, 4, 0, 0, 4, 8, 0, 2, 16, 32, 0, 2, 2, 16)^T$$

To calculate the q-value, we will use our state vector acting on the weight vector representing each action, which can be written as:

$$q_a = \phi^T w_a$$

where  $q_a$  is the q value of that action and  $w_a$  is the weight of that action.

		2	4
		4	8
	2	16	32
	2	2	16

Figure 2: An example game board. This image is obtained from Wikipedia.

### 3.2.2 Relationship Representation

Reading Barto's "Reinforcement Learning" book, I have learnt that, to improve my learning, my state vector needs to have information involving interaction between the tiles. In addition to that, I have also decided to use the book's convention where the action dependency will be put in the state instead of the weight vector. This means our Linear Approximation equation can be re-written as:

$$q_a = \phi(a)^T w$$

When we play the game, whenever we consider a horizontal move, we always compare the tiles horizontally. This means the interaction considered will be horizontal. We will represent the interaction by the difference between 2 tiles. Our state vector for move left will then be written as:

$$\phi(L) = (\text{values of all 16 tiles, horizontal differences between each tile of each row})^T$$

To distinguish move "L" and "R" the horizontal differences in  $\phi(R)$  will have opposite sign to that in  $\phi(L)$ .  $\phi(U)$  and  $\phi(D)$  can be calculated with similar method except we will consider vertical differences instead.

### 3.3 Actions

As explained before, there are 4 available actions for the agent to choose from. They are "Left", "Right", "Up" and "Down". Our agent will choose the action by calculating the q-value of each action. Since there are lots of possible traces in this game, we choose the action using the epsilon-greedy policy to allow exploration by the agent. The epsilon-greedy policy is:

$$\pi(s, a) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|A|}, & a = \text{argmax}_{a'}(q(s, a')) \\ \frac{\epsilon}{|A|} & \text{otherwise} \end{cases} \quad (3)$$

where  $|A|$  is the number of available actions the agent has. In our case,  $|A|$  will be most likely to be 4. As mentioned above, there are a huge possibility of traces in this game. To ensure we don't waste time exploring invalid moves, we have decided to restrict our agent to only choose from valid moves. This means if "Left" is not a valid move, we will only have 3 actions available for the agent to choose from and  $|A| = 3$  in this case instead assuming the other 3 moves are still valid.

From multiple sources, I have also seen that instead of using a pure epsilon-greedy policy, people tend to use a decaying epsilon-greedy policy. This allows exploration at the beginning but it decreases the freedom of exploration as we get to the later episodes of the learning process. We have implemented this by a rather naive step-wise decaying scheme.

### 3.4 Reward Schemes

Multiple reward schemes has been attempted and they both have their merit and drawbacks. We started off with a simple reward scheme where it will reward the value of the greatest value on the board after a move.

For example, if after a move we obtain a board state shown in figure 2, we will get 32 points.

Another reward scheme we have attempted will be the merge-reward scheme, where the agent will be rewarded whenever we merge two tiles into one. The reward will be the value of the new tile. For example, if the initial board is that shown in figure 2 and we move right, the two "2"s in the bottom row will merge and give a new tile "4". This will give us 4 points.

### 3.5 Learning

We have mainly used the q-learning equation when updating the weight vector, which is:

$$\begin{aligned}\Delta w_a &= \alpha(R_{t+1} + \gamma \max_a(S_{t+1}(a), w_{t+1}) - q(S_t; w_t) \nabla_w q(S_t; w_t)) \\ &= \alpha(R_{t+1} + \gamma \max_a(S_{t+1}(a), w_{t+1}) - q(S_t; w_t) S_t)\end{aligned}\tag{4}$$

where we have taken  $\alpha = 0.0000001$ ,  $\gamma = 0.9$ .

We used the decaying-epsilon-greedy policy as explained above with  $\epsilon = 0.1$  initially. Finally, we have decided to play the game with at least 1000 episodes.

## 4 Results

Your results section text goes here...

## 5 Discussion

Your discussion section text goes here...

## Declaration

This document represents the group report submission from the named authors for the project assignment of module: Foundations of Machine Learning and Data Science (INST0060), 2018-19. In submitting this document, the authors certify that all submissions for this project are a fair representation of their own work and satisfy the UCL regulations on plagiarism.

## References

- [1] Christopher M. Bishop. *Pattern Recognition And Machine Learning*. Springer-Verlag New York, 2006.
- [2] Tom Fawcett. An Introduction to ROC Analysis. *Pattern Recognition Letters*, 27:861–874, 2006.
- [3] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Distributed Representations of Words and Phrases and their Compositionality. In *Advances in Neural Information Processing Systems (NIPS)*, pages 3111–3119, 2013.